

Ten Simple Rules to increase computational skills among biologists with Code Clubs

Ada K. Hagan^{1,2}, Nicholas A. Lesniak¹, Marcy J. Balunas³, Lucas Bishop¹, William L. Close¹, Matthew D. Doherty¹, Amanda G. Elmore¹, Kaitlin J. Flynn^{1,4}, Geoffrey D. Hannigan^{1,5}, Charlie C. Koumpouras^{1,6}, Matthew L. Jenior^{1,7}, Ariangela J. Kozik⁸, Kathryn McBride¹, Samara B. Rifkin⁸, Joshua M.A. Stough¹, Kelly L. Sovacool⁹, Marc A. Sze^{1,10}, Sarah Tomkovich¹, Begum D. Topcuoglu^{1,5}, and Patrick D. Schloss¹

To whom correspondence should be addressed: pschloss@umich.edu

1. Department of Microbiology and Immunology, University of Michigan, Ann Arbor, MI
2. Current affiliation: Alliance SciComm & Consulting, Linden, MI
3. Division of Medicinal Chemistry, Department of Pharmaceutical Sciences, University of Connecticut, Storrs, CT
4. Current affiliation: Benaroya Research Institute, Seattle, WA
5. Current affiliation: Exploratory Science Center, Merck & Co., Inc., Cambridge, Massachusetts, USA
6. Current affiliation: Roche Diagnostics, Clinical Operations Services & eSystems, Indianapolis, IN
7. Current affiliation: Department of Biomedical Engineering, University of Virginia, Charlottesville, VA
8. Department of Internal Medicine, University of Michigan
9. Department of Computational Medicine & Bioinformatics, University of Michigan
10. Current affiliation: Genomics and Pharmacogenomics, Merck & Co., Inc., Cambridge, MA

Ten Simple Rules

1 Introduction

2 For most biologists, the ability to generate data has outpaced the ability to analyze those data.
3 High throughput data comes to us from DNA and RNA sequencing, flow cytometry, metabolomics,
4 molecular screens, and more. Although some accept the approach of compartmentalizing data
5 generation and data analysis, we have found scientists feel empowered when they can both ask
6 and answer their own biological questions [1]. In our experience performing microbiome research,
7 it is more common to find exceptional bench scientists who are inexperienced at analyzing large
8 data sets than to find the reverse. Of course this raises a challenge: How do we train bench
9 scientists to effectively answer biological questions with these larger datasets?

10 The standard undergraduate and graduate training in the biological sciences is typically insufficient
11 for developing these skills [2]. To fill this void, there has been growth in online data science
12 training resources through companies including Data Camp (<https://www.datacamp.com>) and
13 Codecademy (<https://www.codecademy.com>), massive open online courses (MOOCs), and in
14 person workshops hosted by organizations including The Carpentries (<https://carpentries.org>).
15 Each of these formats are highly popular among their audience (e.g. [3]); however, recent
16 analyses of workshops suggests that they have little impact on the long term retention of the
17 material that is taught [4]. This result should not be surprising. These learning formats ask
18 participants to engage in massed practice; an approach that is known to be ineffective [5–8]. In
19 contrast, participants need regular opportunities to engage in repeated deliberate practice that
20 encourage them to practice their retrieval and application skills of the material [9,10]. There is
21 clearly a need to develop a framework that can build upon material introduced in coursework
22 and workshops and that can help researchers to stay up to date on the latest best practices,
23 algorithms, and tools.

24 To address this need, we noted the similarity between the feeling of unease in analyzing the data
25 and the struggles scientists face with engaging the voluminous scientific literature. A common
26 strategy for keeping up with the literature is by participating in a journal club. Regardless of
27 the discipline, a Journal Club involves group discussion of a pre-selected paper that range from
28 informal discussions to PowerPoint presentations and course credit. In addition to building upon

material from traditional coursework and staying current on the literature, journal clubs help strengthen skills in critical thinking, communication, and integrating the literature [11]. Since most journal clubs occur on a regular schedule, they are effective by virtue of repeated practice. With this model in mind, over the past four years we have experimented with creating a Code Club model with the goal of improving reproducible data analysis skills in a community environment.

Our Code Club sessions are an hour long and alternate with our lab's Journal Club as the second part of weekly two hour lab meetings. Initially, the Code Club was used to review code from trainee projects. Instead of a presentation, the presenter would project their code onto a screen and the participants would go through the code, stating the logic behind each line. This approach emphasized the importance of code readability and gave beginners the opportunity to see the real-life, messy code of more experienced peers. Unfortunately, the format only allowed us to review a fraction of a project's code, making it difficult to integrate the programmer's logic across their full project. A major issue with this model was that sometimes beginners could not contribute to improving the code, and even when they could, more experienced group members would eventually dominate the discussion. This led to a lack of participation by beginners, who would sometimes mentally withdraw, resulting in an adversarial environment between the presenter and more experienced members. As a result, presenters were reluctant to offer to present again.

From these experiences, we began a collective conversation to improve our Code Club model. We have identified two successful approaches. The first is a more constructive version of a group code review. The presenter clearly states the problem they want to solve, breaks the participants into smaller groups, and then asks each group to solve the problem, or a portion of it. For example, someone may have an R script with a repeating chunk of code. The challenge for the session would be to convert the code chunk into a function to be called throughout the script to make it "DRY" (i.e. Don't Repeat Yourself [12]). The presenter leaves the session with several partial or working solutions to their problem and the importance of writing DRY code is reinforced. The second approach is a tutorial. The presenter introduces a new package or technique and assigns an activity to practice the new approach. For example, at one Code Club participants were given raw data and a finished plot. Paired participants were tasked with generating the plot from the data using R syntax from either the base language or the ggplot2 package. For this exercise, base

R users had to use ggplot2 and vice versa. In either approach, the Code Club ends with a report back to the larger group describing the approach each pair took. We generally find that preparing a Code Club session takes similar effort to preparing a journal club presentation. Our Code Clubs typically have 7 to 10 participants, but the inherent “think-pair-share” approach should allow it to be scaled to larger groups [13].

We continue to experiment with approaches for running Code Club. For example, during the early period of the Covid-19 pandemic in 2020 when many research labs were shuttered, PDS started posting Code Club materials (https://www.riffomonas.org/code_club/) that individuals or research groups can use to practice their coding skills. During this time, we continued our group’s Code Club sessions using the breakout room and screen sharing features in Zoom. Regardless of our approach to leading a Code Club, we have learned that it is critical for the presenter to clearly articulate their goals and facilitate participant engagement. Although some Code Club sessions may be more experimental than others, on the whole they are a critical tool to train bench scientists in reproducible data analysis practices. We have provided some examples of successful topics in Table 1. We have summarized the results of our own experimentation as Code Club presenters and participants into Ten Rules. The first 3 rules apply to all participants, Rules 4 through 8 are targeted to presenters, and the last two for non-presenters. As the rules describe, this model can easily be adapted to groups of non-biologists or those with stronger backgrounds in computer science.

Rule 1: Reciprocate respect

It is critical that the presenter and participants respect each other and that a designated individual (e.g. the lab director) enforces a code of conduct. Each member of the Code Club must have the humility to acknowledge that they have more to learn about any given topic. We have found that many problems are avoided when the presenter takes charge of the session with a clear lesson plan, thoughtfully creates groups, and gives encouragement. Similarly, participants foster a positive environment by remembering that the task is not a competition, instead focusing on the presenter’s goals, allowing their partner to contribute, asking clarifying questions when

85 appropriate, and avoiding distractions (e.g. email, social media). Learning to program is
86 challenging. Too often it is attempted in a environment with nonconstructive criticism. All parties
87 in a Code Club are responsible for preventing this by demonstrating respect for themselves and
88 their colleagues.

89 **Rule 2: Let the material change you**

90 Part of the humility required to participate in a Code Club is acknowledging that your training is
91 incomplete and that it is possible for everyone to learn something new. For participants, assume
92 that the presenter has a plan and follow their presented approach. After the Code Club, try to
93 incorporate that material into new code or refactor old code. By practicing the material in a different
94 context, you will learn the material better. For presenters, incorporate suggested changes into
95 your code. Either party may identify concepts that they are unsure of, presenting opportunities for
96 further conversation and learning.

97 **Rule 3: Experiment!**

98 The selection of content and structure for each Code Club works best when it is democratic and
99 distributed. If someone thinks a technique is worth learning and wants to teach it, they have that
100 power. If they want to experiment with a different format, they are free to try it out. Members
101 of the Code Club need to feel like they have the power to shape the direction of the group. If
102 members are following these Rules, they will naturally reflect on the skills and interests of the
103 other members in the group. For example, there is always turnover in a research group, making it
104 important to revisit basic concepts to teach to new people and provide a refresher to others. One
105 successful experiment found that we could keep basic content interesting to more experienced
106 group members by creating related problems that varied in their difficulty. The group should also
107 feel free to experiment with the format and incorporate group feedback by ending with a debrief
108 discussing the pros and cons of new formats.

Rule 4: Set specific goals

In our early Code Clubs we noticed that if the presenter did not clearly state their goals for the session, it often led to frustration for both the presenter and participants. If the presenter shared their own code, did they want participants to focus on their coding style or did they want help incorporating a new package into their workflow? Participants will always notice or ask about code concepts that are not the focus of the exercise; a presenter with a specific goal can bring tangential conversations back to the planned task. Where possible, presenters should create a simplified scenario (i.e. minimal, reproducible example [14]), which can be helpful in focusing the participants. The presenter should verify ahead of time that the simplified example works and behaves the way they expect. Beyond the content, clear goals for participant activities will help both parties stay on task and avoid frustration. For more advanced learners, the presenter can create stretch goals or give an activity with multiple stopping points where participants would feel successful (e.g. commenting code, creating function, implementing function, refactoring function). Accomplishing specific goals is more likely to result in a positive outcome for presenters and participants.

Rule 5: Keep it simple

Our Code Club needs to fit within an hour time slot. When considering their Code Club activity, the presenter should plan for an introduction and brief instruction, time for participants to engage the material, and time for everyone to report back within that hour. A typical schedule for Code Club is 10 minutes of introduction and instruction, 30 minutes of paired programming, 5 minutes to get groups to wrap up, and 10 minutes to report back to the group. We once had a presenter try to teach basic, but unfamiliar, Julia syntax. Unfortunately, the time was up before the participants had installed the interpreter. Some tips to help keep it simple are to limit the presented code chunks to less than 50 lines or, conversely, consider the number of lines that might be required to accomplish a solution. Remember that learners may need up to three times as long to complete a task that is straight forward for the presenter, so Code Club is best kept simple.

Rule 6: Give participants time to prepare

Similar to a Journal Club, the presenter should give participants a few days (ideally a week) to prepare for the Code Club. Considering the compressed schedule described in Rule 5, asking participants to download materials beforehand is helpful to ensure a quick start. The presenter should provide the participants with instructions on how to install dependencies, download data, and get the initial code. This might also uncover weak points in the presenter's plan and enable them to ensure that the materials work as intended before the Code Club. We have found that using GitHub repositories for each Code Club can help make information, scripts, and data easily available to participants. Although this can be convenient, introducing GitHub on top of the session's activities can impose a significant cognitive load and frustration to those not already comfortable with GitHub. Perhaps a first Code Club could introduce using git and GitHub to engage in collaborative coding. A lower barrier entry point for the presenter is posting their code, data, and information in a lab meeting-dedicated Slack channel or via email. Whatever method is used, the presenter should be sure to communicate the topic and necessary materials with the participants ahead of time.

Rule 7: Don't give participants busy work

Participants want to learn topics that will either be useful to them or help their colleague (i.e. the presenter). Presenters should do their best to satisfy those motivations, whether it is through the relevancy of the concept or the data. It does not make sense to present a Code Club on downloading stock market data if it is not useful or interesting to the group. Similarly, participants should not be tasked with improving the presenter's code if the presenter has no intention of incorporating the suggestions. A list of packages or tasks that group members are interested in along with a log of previous topics could help a presenter struggling to find a topic choose one that will make a rewarding Code Club.

Rule 8: Include all levels of participants

As suggested by Rule 7, a significant challenge to presenting at Code Club is selecting topics and activities that appeal to a critical mass of the participants. This is particularly difficult if participants have a wide range of coding experience, which can follow a turnover in group membership. Beginners will benefit from sessions that cover fundamental concepts and functions. The benefits for more experienced participants include the improved understanding of concepts by teaching and breaking down problems into simpler elements. Furthermore, instead of focusing on core functions, the group could balance Code Club sessions that cover basics with those that introduce new methods and packages to the group. We have also identified several strategies to overcome the challenges presented by participants at varying skill levels. Central to the Code Club format is the use of paired-programming [15]. Instead of letting participants form their own pairs, the presenter can select pairs of participants with either similar or differing skill levels, depending on the presenter's goals. Partnerships between those with similar skill levels requires the presenter to design appropriate activities for each skill level. We have found that commenting code is a good skill for beginners since it forces them to dissect and understand a code chunk line-by-line. It also reinforces the value of commenting as they develop their skills and independence. An advantage of forming partnership between people with disparate skill levels is that it is more likely for groups to provide the presenter with a diverse range of methods that achieve the same result. This approach to pairing also helps to graft new members that have emerging programming skills. Regardless of how partners are selected, consider asking the pairs to identify a navigator and a driver [16]. The driver types at the computer while the navigator tells them what to type, thus ensuring participation of both partners. Midway through the activity, the presenter can have the partners switch roles. Intentionally forming pairs can also engineer group interactions by avoiding potentially disruptive partnerships or pairing reliable role models with new group members.

Rule 9: Prepare in advance to maximize participation

It is not possible to fully participate in a Journal Club discussion about a paper that the participant has not read. In that context, coming to Code Club without having installed a dependency is similar to asking a simple question about the Journal Club paper without first reading it. Both instances show a lack of preparation. Just as a presenter must follow Rule 6 to provide materials ahead of time, participants must review the code in advance, download the data sets, install the necessary packages, and perhaps read up on the topic. If the Code Club is based on a paper or chapters in a data science book (e.g., [17–19]), the participants should read them before the session and consider how they might incorporate the concepts into their own work.

Rule 10: Participate

An essential ingredient of any Code Club is active participation from all parties. Having an open laptop on the table and permission to use it can feel like an invitation to get distracted by other work, emails, and browsing the internet; fight that urge and focus on the presenter's goals. Be respectful and allow your partner to contribute. Speak up for yourself and force your partner to let you contribute. If the material seems too advanced for you, it can be frustrating, and tempting to mentally check out. Fortunately, programming languages like R and Python are generally expressive, which should allow you to engage with the logic, even if the syntax is too advanced. Oftentimes, understanding the logic of *when* to use one modeling approach over another is more important than knowing *how* to use it. If you understand the “why”, the “how” will quickly follow. More experienced participants should aim to communicate feedback and coding suggestions at a level that all participants can understand and engage in. Regardless of skill levels, your partner and the presenter put themselves in a vulnerable position by revealing what they do or do not know. Encourage them, and show your gratitude for helping you learn something new, by fully participating in each Code Club.

207 **Conclusion**

208 The most important rules are the first and last. Members of the Code Club need to feel comfortable
209 with other group members and sufficiently empowered to try something new or ask for help. Aside
210 from expanding our programming skills, we have noticed two other benefits that help create a
211 positive culture. First, we have intentionally interviewed postdoc candidates on the days we hold
212 Code Club. We make it clear that they are not being assessed on their coding skills, but instead
213 use it as an opportunity to see how the candidate interacts with other members of the research
214 group. At the same time, the candidate can learn about the culture of the research group through
215 active participation. Second, members of other research groups have integrated themselves into
216 our Code Club to minimize the isolation they feel in growing their skills within smaller research
217 groups. This speaks to both the broader need for Code Club and the likelihood of success when
218 expanded to include a larger group of individuals with broader research interests. There is no
219 reason that the Code Club format would not work with groups independent of a research group as
220 long as they follow the Rules. Finding datasets and applications that are interesting to a critical
221 mass of people will likely be key to starting and sustaining such a group. Ultimately, Code Club
222 has improved the overall data analysis skills, community, and research success of our lab by
223 empowering researchers to seek help from their colleagues.

224 **Acknowledgements**

225 All co-authors have been participants in the Code Clubs run as part of PDS's lab meetings. They
226 have had a critical role in shaping the evolution of the sessions. AKH led a lab discussion of Code
227 Clubs to identify the key concepts presented, which were supplemented by contributions from
228 NAL. The order of other co-authors is listed alphabetically.

References

1. Carey MA, Papin JA. Ten simple rules for biologists learning to program. Markel S, editor. PLOS Computational Biology. 2018;14: e1005871. doi:10.1371/journal.pcbi.1005871
2. Barone L, Williams J, Micklos D. Unmet needs for analyzing biological big data: A survey of 704 NSF principal investigators. Ouellette F, editor. PLOS Computational Biology. 2017;13: e1005755. doi:10.1371/journal.pcbi.1005755
3. Wilson G. Software carpentry: Lessons learned. F1000Research. 2016. doi:10.12688/f1000research.3-62.v2
4. Feldon DF, Jeong S, Peugh J, Roksa J, Maahs-Fladung C, Shenoy A, et al. Null effects of boot camps and short-format training for PhD students in life sciences. Proceedings of the National Academy of Sciences. 2017;114: 9854–9858. doi:10.1073/pnas.1705783114
5. Carpenter SK, Cepeda NJ, Rohrer D, Kang SHK, Pashler H. Using spacing to enhance diverse forms of learning: Review of recent research and implications for instruction. Educational Psychology Review. 2012;24: 369–378. doi:10.1007/s10648-012-9205-z
6. Rohrer D. Student instruction should be distributed over long time periods. Educational Psychology Review. 2015;27: 635–643. doi:10.1007/s10648-015-9332-4
7. Budé L, Imbos T, Wiel MW van de, Berger MP. The effect of distributed practice on students' conceptual understanding of statistics. Higher Education. 2010;62: 69–79. doi:10.1007/s10734-010-9366-y
8. Karpicke JD, Blunt JR. Retrieval practice produces more learning than elaborative studying with concept mapping. Science. 2011;331: 772–775. doi:10.1126/science.1199327
9. Rohrer D, Dedrick RF, Stershic S. Interleaved practice improves mathematics learning. Journal of Educational Psychology. 2015;107: 900–908. doi:10.1037/edu0000001
10. Karpicke JD, Roediger HL. The critical importance of retrieval for learning. Science. 2008;319: 966–968. doi:10.1126/science.1152408

- 253 11. Lonsdale A, Penington JS, Rice T, Walker M, Dashnow H. Ten simple rules for a
254 bioinformatics journal club. Lewitter F, editor. PLOS Computational Biology. 2016;12: e1004526.
255 doi:10.1371/journal.pcbi.1004526
- 256 12. Wilson G, Aruliah DA, Brown CT, Hong NPC, Davis M, Guy RT, et al. Best practices for
257 scientific computing. Eisen JA, editor. PLoS Biology. 2014;12: e1001745. doi:10.1371/journal.pbio.1001745
- 258 13. Kothiyal A, Murthy S, Iyer S. Think-pair-share in a large CS1 class. Proceedings of the 2014
259 conference on innovation & technology in computer science education - ITiCSE 14. ACM Press;
260 2014. doi:10.1145/2591708.2591739
- 261 14. Stack Overflow. How to create a minimal, reproducible example. 2020. Available: <https://stackoverflow.com/help/minimal-reproducible-example>
262
- 263 15. Hanks B, Fitzgerald S, McCauley R, Murphy L, Zander C. Pair programming in education: A
264 literature review. Computer Science Education. 2011;21: 135–173. doi:10.1080/08993408.2011.579808
- 265 16. Böckeler B, Siessegger N. On pair programming. 2020. Available: <https://martinfowler.com/articles/on-pair-programming.html>
266
- 267 17. Wickham H, Grolemund G. O'Reilly Media; 2016. Available: <https://r4ds.had.co.nz>
- 268 18. Schloss PD. The Riffomonas reproducible research tutorial series. Journal of Open Source
269 Education. 2018;1: 13. doi:10.21105/jose.00013
- 270 19. Noble WS. A quick guide to organizing computational biology projects. Lewitter F, editor. PLoS
271 Computational Biology. 2009;5: e1000424. doi:10.1371/journal.pcbi.1000424

272 **Table 1. Examples of successful Code Club topics.**

| Title | Description |
|-----------------------|---|
| base vs. ggplot2 | Given input data and a figure, recreate the figure using R's base graphics or ggplot2 syntax |
| Snakemake | Given a bash script that contains an analysis pipeline, convert it to a Snakemake workflow (https://github.com/SchlossLab/snakemake_riffomonas_tutorial) |
| DRYing code | Given script with repeated code, create functions to remove repetition |
| Vegan R package | Compare microbial communities using the <code>adonis</code> function in the Vegan R package (https://github.com/SchlossLab/Code_Review_42717) |
| tidy data | Given a wide-formatted data table, convert it to a long, tidy-formatted data table using tools from R's tidyverse |
| GitFlow | Participants file and claim an issue to add their name to a README file in a GitHub-hosted repository and file a pull request to complete the issue |
| googledocs4 R package | Scrape a Google docs workbook and clean the data to identify previous Code Club presenters |
| Develop an R Package | Convert a lab member's collection of scripts into an R package over a series of sessions |
| Documenting R code | Use <code>roxygen2</code> to supplement comments in R code to improve documentation (https://github.com/SchlossLab/documenting-R) |
| gganimate R package | Convert static plots generated with ggplot2 package into GIFs (https://github.com/SchlossLab/2020-04-12-CodeClub_PlotAnimation/) |